

## Concurrent Queues and Stacks

Companion slides for  
The Art of Multiprocessor  
Programming  
by Maurice Herlihy & Nir Shavit

## The Five-Fold Path

- Coarse-grained locking
- Fine-grained locking
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

2

We saw 4 approaches to concurrent data structure design.

## Another Fundamental Problem

- We told you about
  - Sets implemented using linked lists
- Next: queues
- Next: stacks

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

3

## Queues & Stacks

- Both: pool of items
- Queue
  - enq() & deq()
  - First-in-first-out (FIFO) order
- Stack
  - push() & pop()
  - Last-in-first-out (LIFO) order

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

4

We have already talked about concurrent lists, and today we look at another ubiquitous data structure: queues. Queues provide enqueue and dequeue methods and they are used to route requests from one place to another. For example, a web site might queue up requests for pages, for disk writes, for credit card validations and so on.

## Bounded vs Unbounded

- **Bounded**
  - Fixed capacity
  - Good when resources an issue
- **Unbounded**
  - Holds any number of objects

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

5

When designing a pool interface, one choice is whether to make the pool bounded or unbounded. A bounded pool has a fixed capacity (maximum number of objects it holds). Bounded pools are useful when resources are an issue. For example, if you are concerned that the producers will get too far ahead of the consumers, you should bound the pool. An unbounded pool, by contrast, holds an unbounded number of objects.

## Blocking vs Non-Blocking

- **Problem cases:**
  - Removing from empty pool
  - Adding to full (bounded) pool
- **Blocking**
  - Caller waits until state changes
- **Non-Blocking**
  - Method throws exception

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

6

A second choice is what to do when a thread tries to remove an item from an empty queue or to add one to a full bounded queue. One choice is to block the thread until the state changes. This choice makes sense when there is nothing else for the thread to do. The alternative is to throw an exception. This choice makes sense when the thread can profitably turn its attention to other activities.

## This Lecture

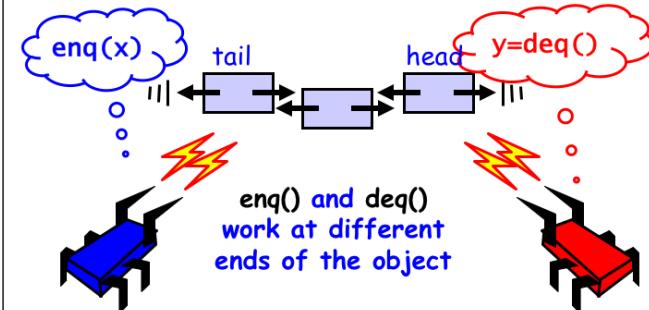
- Bounded, Blocking, Lock-based Queue
- Unbounded, Non-Blocking, Lock-free Queue
- ABA problem
- Unbounded Non-Blocking Lock-free Stack
- Elimination-Backoff Stack

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

7

In this lecture we are going to look at two alternative highly-concurrent queue implementations. The first is a bounded, blocking, lock-based queue. This implementation uses a combination of techniques due to Doug Lea and to Maged Michael and Michael Scott. The second implementation is an unbounded, non-blocking, lock-free implementation due to Michael and Scott. We will then continue to examine the design of concurrent stacks, show Treiber's lock-free stack and then how one can add parallelism to a stack through a technique called elimination.

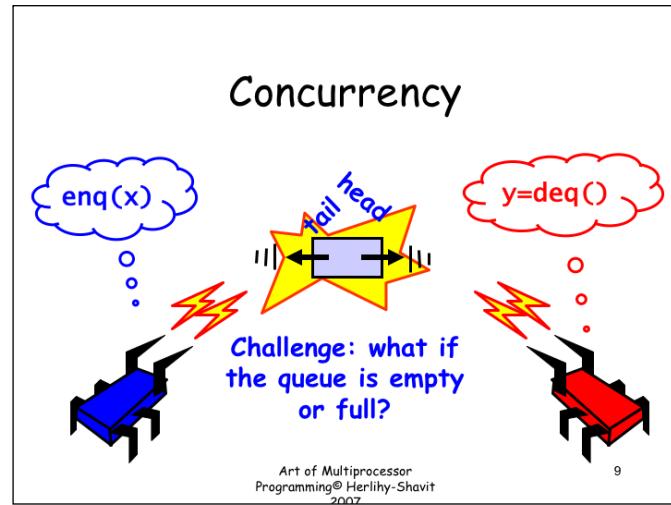
## Queue: Concurrency



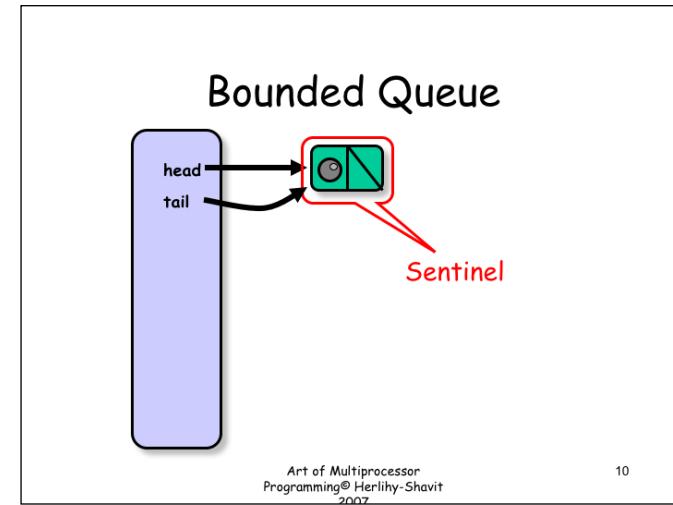
Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

8

Making a queue concurrent is quite a challenge. Very informally, it seems that it should be OK for one thread to enqueue an item at one end of the queue while another thread dequeues an item from the other end, since they are working on disjoint parts of the data structure. Does that kind of concurrency seem easy to realize? (Answer: next slide)

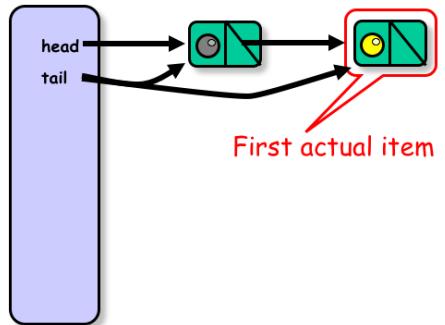


Concurrent enqueue and dequeue calls should in principle be able to proceed without interference, but it gets tricky when the queue is nearly full or empty, because then one method call should affect the other. The problem is that whether the two method calls interfere (that is, need to synchronize) depends on the dynamic state of the queue. In pretty much all of the other kinds of synchronization we considered, whether or not method calls needed to synchronize was determined statically (say read/write locks) or was unlikely to change very often (resizable hash tables).



Let's use a list-based structure, although arrays would also work. We start out with head and tail fields that point to the first and last entries in the list. The first Node in the list is a sentinel Node whose value field is meaningless. The sentinel acts as a placeholder.

## Bounded Queue

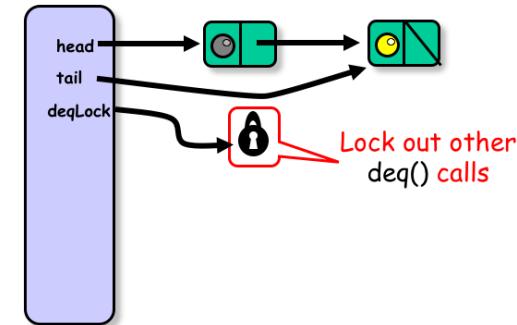


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

11

When we add actual items to the queue, we will hang them off the sentinel Node at the head.

## Bounded Queue

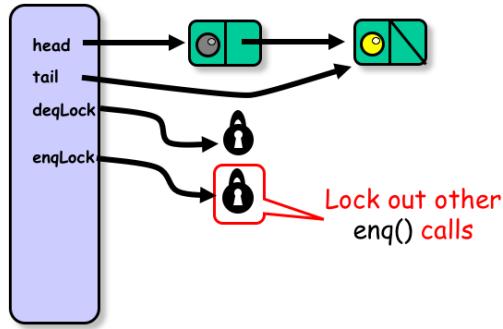


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

12

The most straightforward way to allow concurrent enq and deq calls is to use one lock at each end of the queue. To make sure that only one dequeuer is accessing the queue at a time, we introduce a dequeue lock field. Could we get the same effect just by locking the tail Node? (Answer: there's a race condition between when a thread reads the tail field and when it acquires the lock. Another thread might have enqueued something in that interval. Of course you could check after you acquire the lock that the tail pointer hadn't changed, but that seems needlessly complicated.)

## Bounded Queue

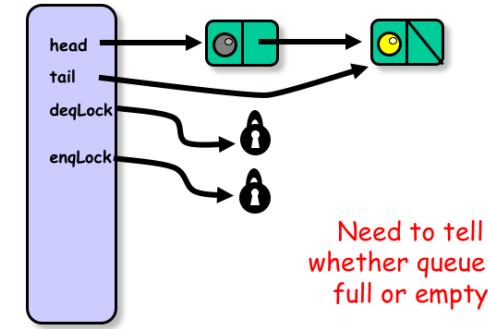


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

13

In the same way, we introduce an explicit `enqLock` field to ensure that only one enqueue can be accessing the queue at a time. Would it work if we replaced the `enqLock` by a lock on the first sentinel Node? (Answer: Yes, it would because the value of the `head` field never changes. Nevertheless we use an `enqLock` field to be symmetric with the `deqLock`.)

## Not Done Yet

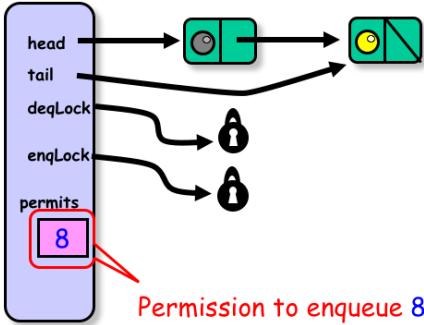


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

14

In the same way, we introduce an explicit `enqLock` field to ensure that only one enqueue can be accessing the queue at a time. Would it work if we replaced the `enqLock` by a lock on the first sentinel Node? (Answer: Yes, it would because the value of the `head` field never changes. Nevertheless we use an `enqLock` field to be symmetric with the `deqLock`.)

## Not Done Yet

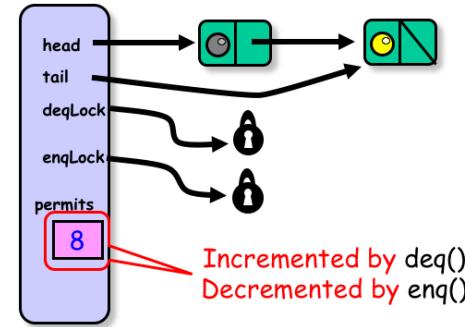


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

15

Let's add another field, which we will think of as keeping track of permissions to enqueue items. We could also think of this field as counting the number of empty slots, but calling it a permission will seem more natural later on when we look at some additional fine-grained tricks.

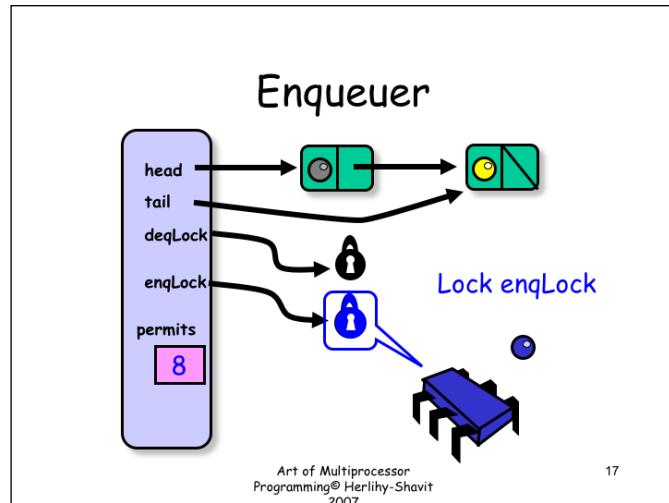
## Not Done Yet



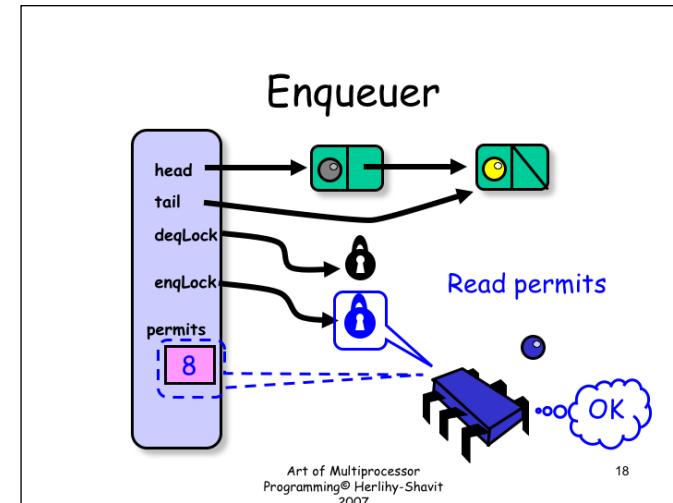
Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

16

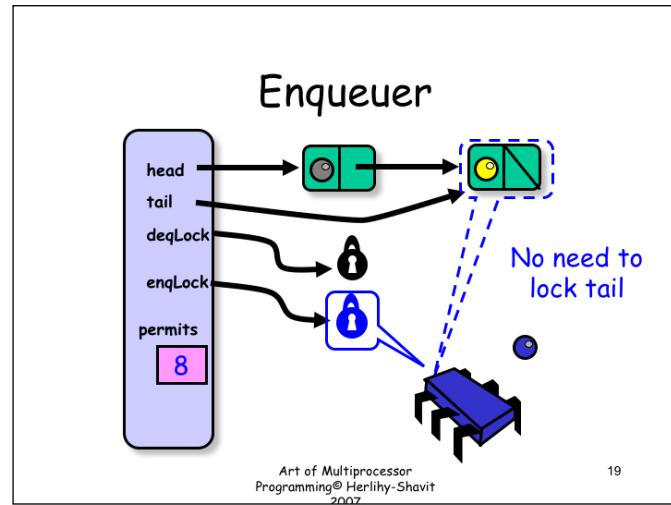
This field is incremented by the `deq` method (which creates a space) and decremented by the `enq` method.



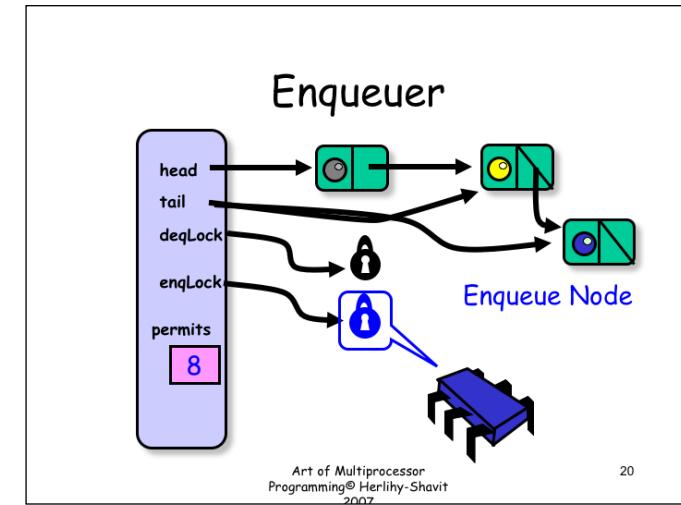
Let's walk through a very simple implementation. A thread that wants to enqueue an item first acquires the enqLock. At this point, we know there are no other enq calls in progress.



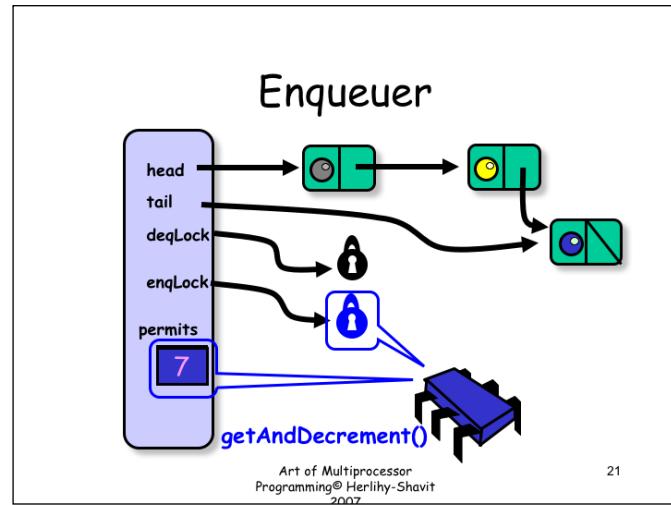
Next, the thread reads the number of permits. What do we know about how this field behaves when the thread is holding the enqLock? (Answer: it can only increase. Dequeueurs can increment the counter, but all the enqueueurs are locked out.) If the enqueueing thread sees a non-zero value for permits, it can proceed.



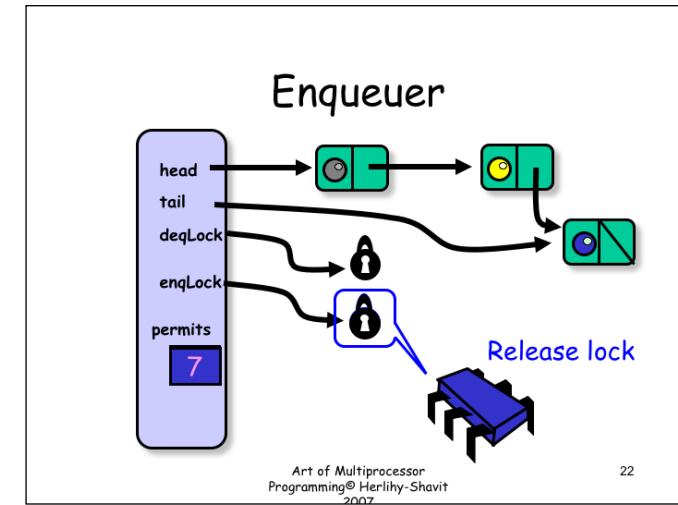
The thread does not need to lock tail Node before appending a new Node. Clearly, there is no conflict with another enqueue and we will see a clever way to avoid conflict with a dequeuer.



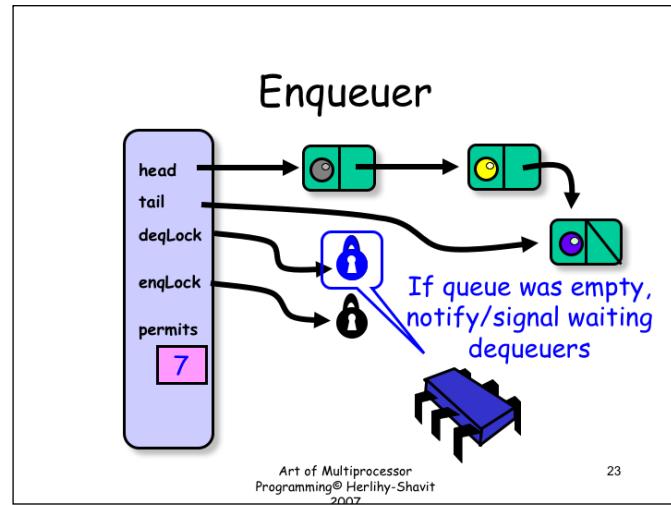
Still holding the enqLock, the thread redirects both the tail Node's next field and the queue's tail field to point to the new Node.



Still holding the enqLock, the thread calls `getAndDecrement()` to reduce the number of permits. Remember that we know the number is positive, although we don't know exactly what it is now. Why can't we release the enqLock before the increment? (Answer: because then other enqueueers won't know when they have run out of permits.)

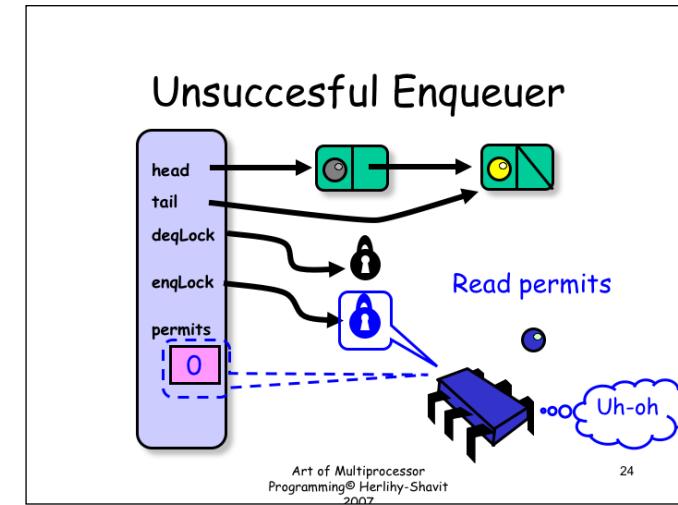


Finally, we can release the enqLock and return.



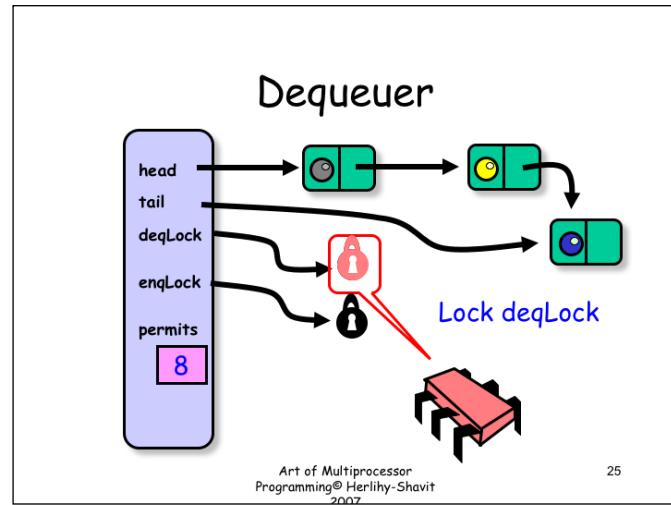
There is one more issue to consider. If the queue was empty, then one or more dequeuers may be waiting for an item to appear. If the dequeuer is spinning then we don't need to do anything, but if the dequeuers are suspended in the usual Java way, then we need to notify them. There are many ways to do this, but for simplicity we will not try to be clever, and simply require any enqueuer who puts an item in an empty queue to acquire the deqLock and notify any waiting threads.

**TALKING POINT:** Notice that Java requires you to acquire the lock for an object before you can notify threads waiting to reacquire that lock. Not clear if this design is a good idea. Does it invite deadlock? Discuss.

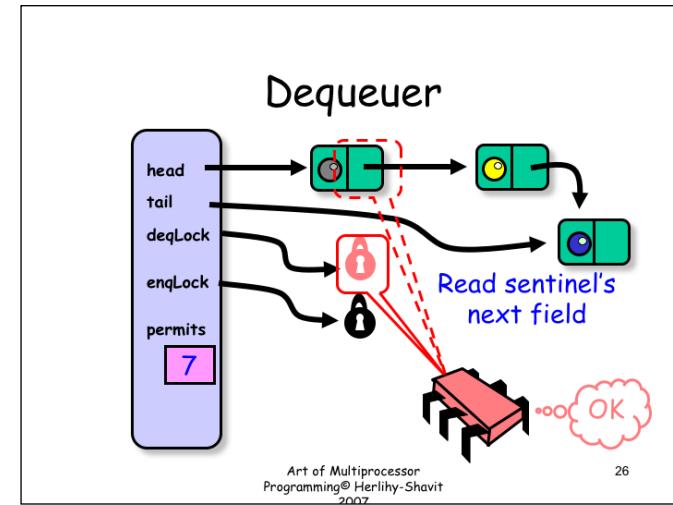


Let's rewind and consider what happens if the enqueueing thread found the queue full. Suppose the threads synchronize by blocking, say using the Java wait() method. In this case, the thread can wait() on the enqLock, provided the first dequeuer to remove an item notifies the waiting thread.

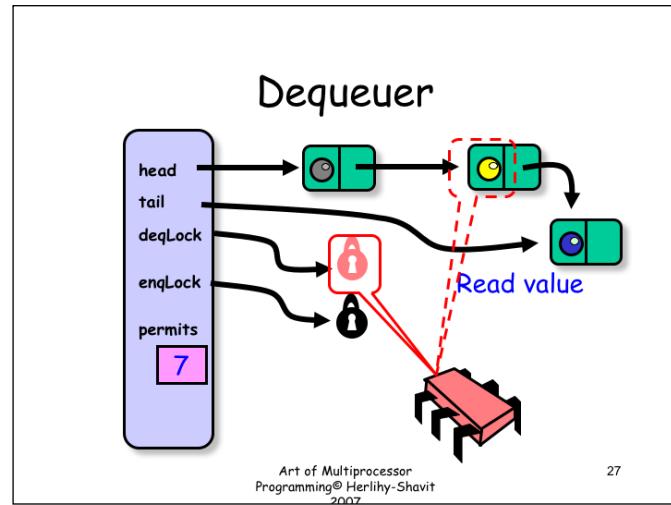
An alternative is for the enqueuer to spin waiting for the number of permits to be non-zero. In general, it is a dangerous practice to spin while holding a lock. It so happens that it works in this case, provided the dequeuer also spins when the queue is empty. Spinning here is cache-friendly, since the thread is spinning on a locally-cached value, and stops spinning as soon as the field value changes.



Now let's walk through a similar `deq()` implementation. The code is similar, but not at all identical. As a first step, the dequeueer thread acquires the `deqLock`.

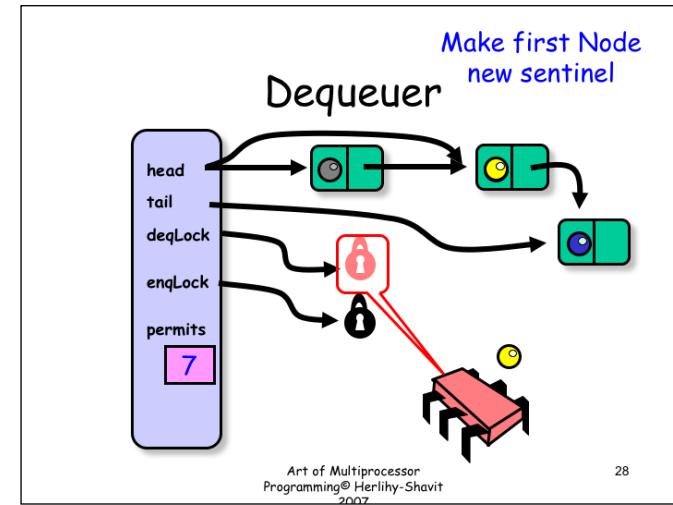


Next, the thread reads the sentinel node's next field. If that field is non-null, the queue is non-empty. Once the enqueueer sees a non-empty next field, that field will remain that way. Why? (ANSWER: the `deq()` code, which we have seen, changes null references to non-null, but never changes a non-null reference). The enqueueers are all locked out, and no dequeueer will modify a non-null field. There is no need to lock the Node itself, because it is implicitly protected by the `enqLock`.



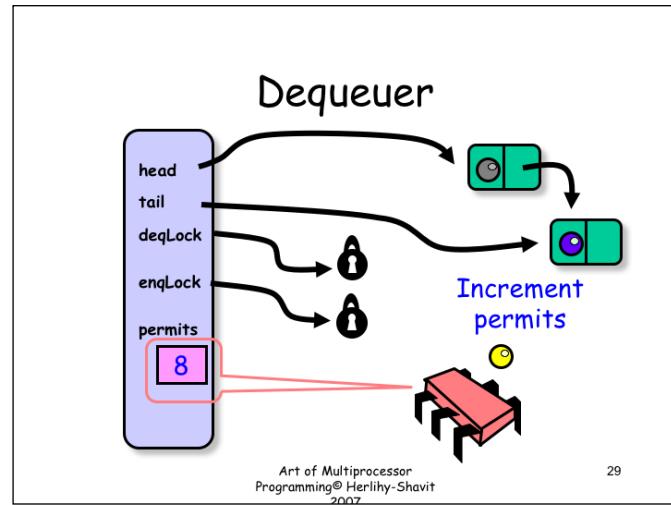
Next, the thread reads the first Node's value, and stores it in a local variable.

27



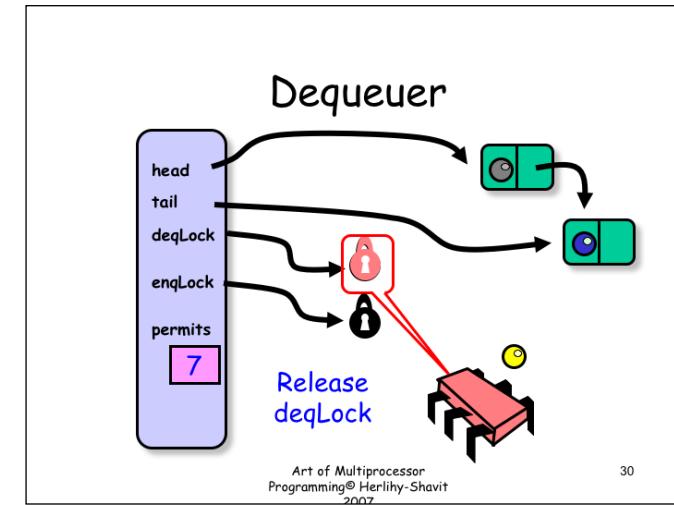
The thread then makes the first Node the new sentinel, and discards the old sentinel. This step may seem obvious, but, in fact, it is enormously clever. If we had tried to physically remove the Node, we would soon become bogged down in quagmire of complications. Instead, we discard the prior sentinel, and transform the prior first real Node into a sentinel. Brilliant.

28



29

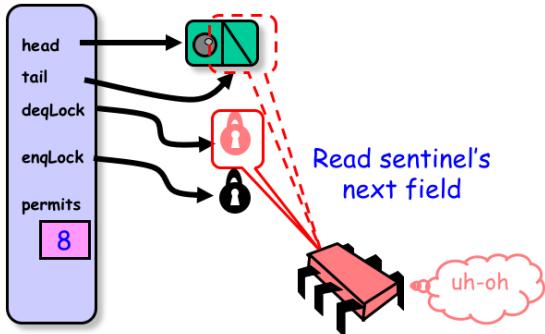
Finally, we increment the number of permits. By contrast with enqueueers, we do not need to hold the lock while we decrement the permits. Why? (Answer: we had to hold the lock while enqueueing to prevent lots of enqueueers from proceeding without noticing that the capacity had been exceeded. Dequeueurs will notice the queue is empty when they observe that the sentinel's next field is null.)



30

Next we can release the deqLo  
ck.

## Unsuccessful Dequeuer



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

31

## Bounded Queue

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger permits;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

32

If the dequeuer observes that the sentinel's next field is null, then it must wait for something to be enqueued. As with the dequeuer, the simplest approach is just to wait() on the deqLock, with the understanding that any enqueueer who makes the queue non-empty will notify any waiting threads. Spinning works too, even though the thread is holding a lock. Clearly, mixing the two strategies will not work: if a dequeuer spins while holding the deq lock and an enqueueer tries to acquire the deq lock to notify waiting threads, then a deadlock ensues.

## Bounded Queue

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger permits;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

Enq & deq locks

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

33

## Digression: Monitor Locks

- Java Synchronized objects and Java ReentrantLocks are monitors
- Allow blocking on a condition rather than spinning
- Threads:
  - acquire and release lock
  - wait on a condition

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

34

In Java, every object provides a `wait()` method that unlocks the object and suspends the caller for a while. While the caller is waiting, another thread can lock and change the object. Later, when the suspended thread resumes, it locks the object again before it returns from the `wait()` call.

## The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock();  
}
```

Acquire lock

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

35

## The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock();  
}
```

Release lock

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

36

## The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock();  
}
```

Try for lock, but not too hard

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

37

## The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock();  
}
```

Create condition to wait on

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

38

## The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock();  
}
```

Guess what this method does?

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

39

## Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

40

Java allows threads to interrupt one another. Interrupts are not preemptive: a thread typically must check whether it has been interrupted. Some blocking methods, like the built-in wait() method throw InterruptedException if the thread is interrupted while waiting. The lock method of the lock interface does not detect interrupts, presumably because it is more efficient not to do so. This variant does.

## Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Release lock and  
wait on condition

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

41

## Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Wake up one waiting thread

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

42

## Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

Wake up all waiting threads

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

43

## Await

q.await()

- Releases lock associated with q
- Sleeps (gives up processor)
- Awakens (resumes running)
- Reacquires lock & returns

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

44

In Java, every object provides a **wait()** method that unlocks the object and suspends the caller for a while. While the caller is waiting, another thread can lock and change the object. Later, when the suspended thread resumes, it locks the object again before it returns from the **wait()** call.

## Signal

```
q.signal();
```

- Awakens one waiting thread
  - Which will reacquire lock

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

45

The **notify()** method (similarly **notify()** for synchronized methods) wakes up one waiting thread, chosen arbitrarily from the set of waiting threads. When that thread awakens, it competes for the lock like any other thread. When that thread reacquires the lock, it returns from its **wait()** call. You cannot control which waiting thread is chosen.

## Signal All

```
q.signalAll();
```

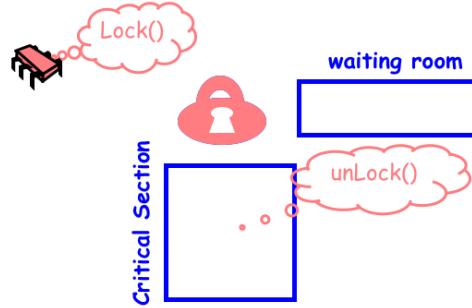
- Awakens all waiting threads
  - Which will each reacquire lock

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

46

The **signalAll()** method wakes up all waiting threads. Each time the object is unlocked, one of these newly-wakened threads will reacquire the lock and return from its **wait()** call. You cannot control the order in which the threads reacquire the lock.

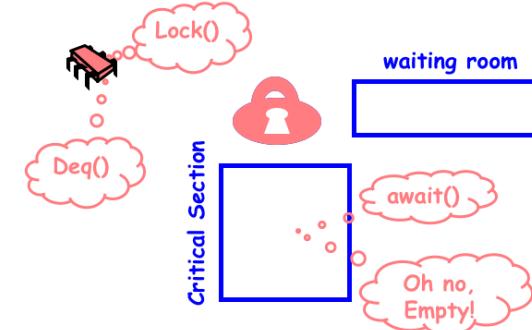
## A Monitor Lock



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

47

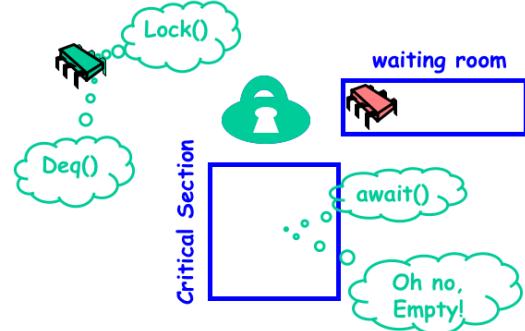
## Unsuccessful Deq



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

48

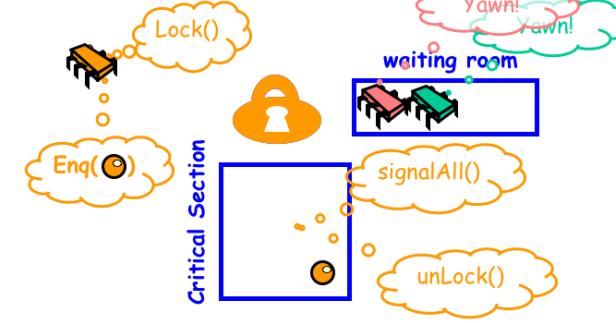
## Another One



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

49

## Enqueur to the Rescue



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

50

## Monitor Signalling



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

51

Awakened thread  
might still lose lock to  
outside contender...

## Dequeurs Signalled



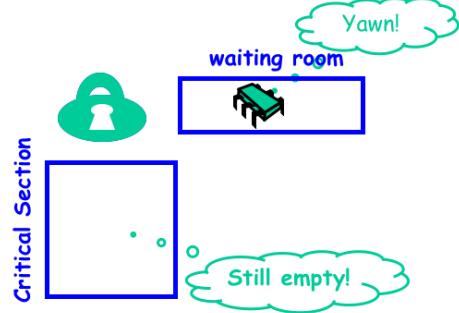
Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

52

Notice, there may be competition from other threads attempting to lock for the thread. The FIFO order is arbitrary, different monitor locks have different ordering of waiting threads, that is, notify could release the earliest or latest any other waiting thread.

Notice, there may be competition from other threads attempting to lock for the thread. The FIFO order is arbitrary, different monitor locks have different ordering of waiting threads, that is, notify could release the earliest or latest any other waiting thread.

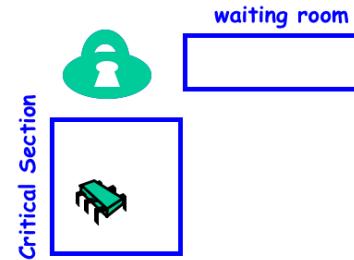
## Dequeurs Signalled



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

53

## Dollar Short + Day Late

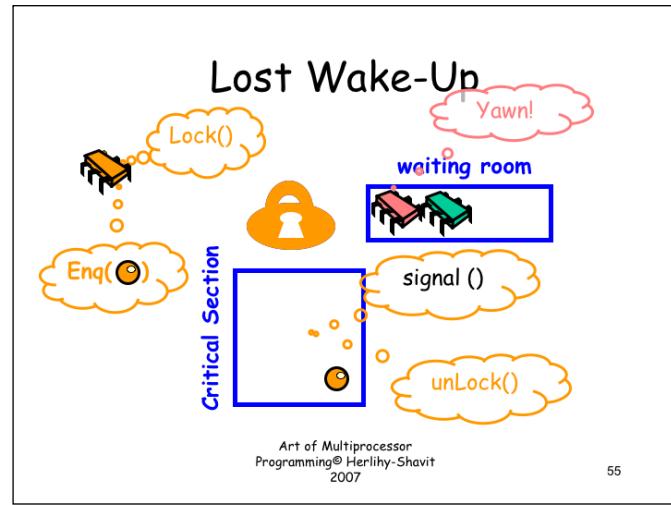


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

54

Notice, there may be competition from other threads attempting to lock for the thread. The FIFO order is arbitrary, different monitor locks have different ordering of waiting threads, that is, notify could release the earliest or latest any other waiting thread.

Notice, there may be competition from other threads attempting to lock for the thread. The FIFO order is arbitrary, different monitor locks have different ordering of waiting threads, that is, notify could release the earliest or latest any other waiting thread.

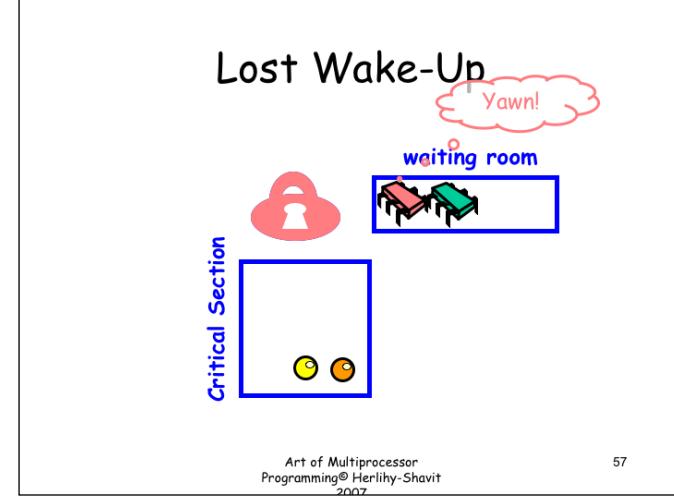
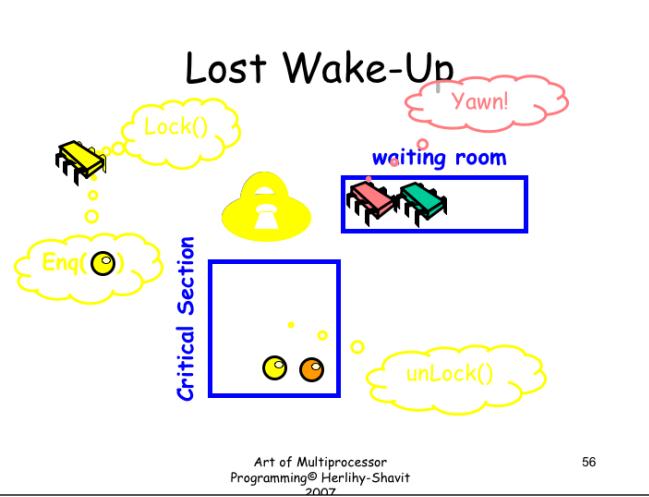


Just as locks are inherently vulnerable to deadlock, `\cCondition` objects are inherently vulnerable to `\emph{lost wakeups}`, in which one or more threads wait forever without realizing that the condition for which they are waiting has become true.

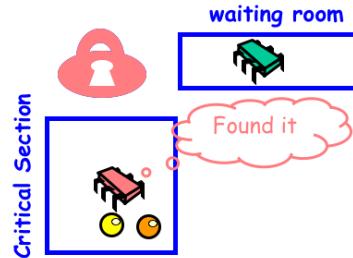
Lost wakeups can occur in subtle ways. Figure `\ref{figure:monitor:lost}` shows an ill-considered optimization of the `\cQueue{T}` class. Instead of signaling the `\fNotEmpty` condition each time `\mEnq` enqueues an item, would it not be more efficient to signal the condition only when the queue actually transitions from empty to non-empty? This optimization works as intended if there is only one producer and one consumer, but it is incorrect if there are multiple producers or consumers. Consider the following scenario:

consumers `$A$` and `$B$` both try to dequeue an item from an empty queue, both detect the queue is empty, and both block on the `\fNotEmpty` condition. Producer `$C$` enqueues an item in the buffer, and signals `\fNotEmpty`, waking `$A$`. Before `$A$` can acquire the lock, however, another producer `$D$` puts a second item in the queue, and because the queue is not empty, it does not signal `\fNotEmpty`. Then `$A$` acquires the lock, removes the first item, but `$B$`, victim of a lost wakeup, waits forever even though there is an item in the buffer to be consumed.

Although there is no substitute for reasoning carefully about your program, there are simple programming practices that will minimize vulnerability to lost wakeups.



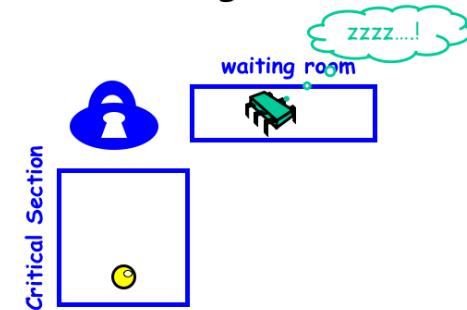
## Lost Wake-Up



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

58

## What's Wrong Here?



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

59

## Solution to Lost Wakeup

- Always use signalAll and notifyAll
- Not signal and notify

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

60

## Java Synchronized Methods

```
public class Queue<T> {  
  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            this.wait();  
        T result = items[head % QSIZE]; head++;  
        this.notifyAll();  
        return result;  
    }  
    ...  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

61

Synchronized methods use a monitor lock also.

## Java Synchronized Methods

```
public class Queue<T> {  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            this.wait();  
        T result = items[head % QSIZE]; head++;  
        this.notifyAll();  
        return result;  
    }  
    ...  
}
```

Each object has an implicit lock with an implicit condition

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

62

## Java Synchronized Methods

```
public class Queue<T> {  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            this.wait();  
        T result = items[head % QSIZE]; head++;  
        this.notifyAll();  
        return result;  
    }  
    ...  
}
```

Lock on entry,  
unlock on return

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

63

Synchronized methods use a monitor lock also.

Synchronized methods use a monitor lock also.

## Java Synchronized Methods

```
public class Queue<T> {  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            this.wait();  
        T result = items[head % QSIZE]; head++;  
        this.notifyAll();  
        return result;  
    }  
    ...  
}
```

Wait on implicit condition

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

64

## Java Synchronized Methods

```
public class Queue<T> {  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            this.wait();  
        T result = items[head % QSIZE]; head++;  
        this.notifyAll();  
        return result;  
    }  
    ...  
}
```

Signal all threads waiting  
on condition

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

65

Synchronized methods use a monitor lock also.

Synchronized methods use a monitor lock also.

## (Pop!) The Bounded Queue

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, dequeLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger permits;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    dequeLock = new ReentrantLock();  
    notEmptyCondition = dequeLock.newCondition();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

66

## Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, dequeLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger permits;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    dequeLock = new ReentrantLock();  
    notEmptyCondition = dequeLock.newCondition();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

Enq & deque locks

67

## Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger permits;  
    Node head; Enq lock's associated  
    condition  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

68

## Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
AtomicInteger permits; Num permits: 0 to capacity  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

69

## Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, dequeLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger permits;  
    Node head; Head and Tail  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    dequeLock = new ReentrantLock();  
    notEmptyCondition = dequeLock.newCondition();  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

70

## Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (permits.get() == 0)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = e;  
        if (permits.getAndDecrement() == capacity)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

71

A remarkable aspect of this queue implementation is that the methods are subtle, but they fit on a single slide.

The `\mEnq()` method (Figure~\ref{figure:boundedQueue:enq}) works as follows.

A thread acquires the `\fEnqLock()` (Line \ref{line:bounded:lock}), and repeatedly reads the `\fPermits[]` field (Line \ref{line:bounded:permits}).

While that field is zero, the queue is full, and the enqueueer must wait until a dequeueuer makes room.

The enqueueer waits by waiting on the `\fNotFullCondition[]` field (Line \ref{line:bounded:notfull}), which releases the enqueue lock temporarily, and blocks until the condition is signaled. Each time the thread awakens (Line \ref{line:bounded:notfull}),

it checks whether the `\fPermits{}` field is positive (that is, the queue is not empty) and if not, goes back to sleep.

## Enq Method Part One

```
public void enq(T x) {
    boolean mustWakeDequeuers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeuers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

*Lock and unlock  
enq lock*

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

72

A remarkable aspect of this queue implementation is that the methods are subtle, but they fit on a single slide.

The `\mEnq{}` method (Figure~\ref{figure:boundedQueue:enq}) works as follows.

A thread acquires the `\fEnqLock{}` (Line \ref{line:bounded:lock}), and repeatedly reads the `\fPermits{}` field (Line \ref{line:bounded:permits}).

While that field is zero, the queue is full, and the enqueueer must wait until a dequeuer makes room.

The enqueueer waits by waiting on the `\fNotFullCondition{}` field (Line \ref{line:bounded:notfull}), which releases the enqueue lock temporarily, and blocks until the condition is signaled. Each time the thread awakens (Line \ref{line:bounded:notfull}),

it checks whether the `\fPermits{}` field is positive (that is, the queue is not empty) and if not, goes back to sleep.

## Enq Method Part One

```
public void enq(T x) {
    boolean mustWakeDequeueurs = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = e;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeueurs = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

If queue is full, patiently await further instructions ...

Art of Multiprocessor Programming® Herlihy-Shavit  
2007

73

Once the number of permits exceeds zero, however, the enqueueer may proceed.

Note that once the enqueueer observes a positive number of permits, then while the enqueue is in progress no other thread can cause the number of permits to fall back to zero, because all the other enqueueers are locked out, and a concurrent dequeuer can only increase the number of permits.

We must check carefully that this implementation does not suffer from a ``lost-wakeup'' bug. Care is needed because an enqueueer encounters a full queue in two steps: first, it sees that `\fPermits{}` is zero, and second, it waits on the `\fNotFullCondition{}` condition until there is room in the queue. When a dequeuer changes the queue from full to not-full, it acquires `\fEnqLock{}` and signals the `\fNotFullCondition{}` condition. Even though the `\fPermits{}` field is not protected by

the `\fEnqLock()`, the dequeuer acquires the `\fEnqLock()` before it signals the condition, so the dequeuer cannot signal between the enqueueer's two steps.

## Be Afraid

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (permits.get() == 0)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = e;  
        if (permits.getAndDecrement() == capacity)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

How do we know the  
permits field won't change?

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

74

Once the number of permits exceeds zero, however, the enqueueer may proceed.

Note that once the enqueueer observes a positive number of permits, then while the enqueue is in progress no other thread can cause the number of permits to fall back to zero, because all the other enqueueers are locked out, and a concurrent dequeuer can only increase the number of permits.

We must check carefully that this implementation does not suffer from a ``lost-wakeup'' bug. Care is needed because an enqueueer encounters a full queue in two steps: first, it sees that `\fPermits()` is zero, and second, it waits on the `\fNotFullCondition()` condition until there is room in the queue. When a dequeuer changes the queue from full to not-full, it acquires `\fEnqLock()` and signals the `\fNotFullCondition()` condition. Even though the `\fPermits()` field is not protected by

the `\fEnqLock{}`, the dequeuer acquires the `\fEnqLock{}` before it signals the condition, so the dequeuer cannot signal between the enqueue's two steps.

## Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (permits.get() == 0)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = e;  
        if (permits.getAndIncrement() == capacity)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

Add new node

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

75

## Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (permits.get() == 0)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = e;  
        if (permits.getAndDecrement() == capacity)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
}
```

If queue was empty, wake  
frustrated dequeuers

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

76

## Enq Method Part Deux

```
public void enq(T x) {  
    ...  
    if (mustWakeDequeueurs) {  
        deqLock.lock();  
        try {  
            notEmptyCondition.signalAll();  
        } finally {  
            deqLock.unlock();  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

77

A remarkable aspect of this queue implementation is that the methods are subtle, but they fit on a single slide.

The `\mEnq()` method (Figure~\ref{figure:boundedQueue:enq}) works as follows.

A thread acquires the `\fEnqLock()` (Line \ref{line:bounded:lock}), and repeatedly reads the `\fPermits[]` field (Line \ref{line:bounded:permits}).

While that field is zero, the queue is full, and the enqueueer must wait until a dequeuer makes room.

The enqueueer waits by waiting on the `\fNotFullCondition[]` field (Line \ref{line:bounded:notfull}), which releases the enqueue lock temporarily, and blocks until the condition is signaled. Each time the thread awakens (Line \ref{line:bounded:notfull}),

it checks whether the `\fPermits{}` field is positive (that is, the queue is not empty) and if not, goes back to sleep.

## Enq Method Part Deux

```
public void enq(T x) {  
    ...  
    if (mustWakeDequeueurs){  
        dequeLock.lock();  
        try {  
            notEmptyCondition.signalAll();  
        } finally {  
            dequeLock.unlock();  
        }  
    }  
}
```

Are there dequeuers to be signaled?

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

78

A remarkable aspect of this queue implementation is that the methods are subtle, but they fit on a single slide.

The `\mEnq{}` method (Figure~\ref{figure:boundedQueue:enq}) works as follows.

A thread acquires the `\fEnqLock{}` (Line \ref{line:bounded:lock}), and repeatedly reads the `\fPermits{}` field (Line \ref{line:bounded:permits}).

While that field is zero, the queue is full, and the enqueueer must wait until a dequeuer makes room.

The enqueueer waits by waiting on the `\fNotFullCondition{}` field (Line \ref{line:bounded:notfull}), which releases the enqueue lock temporarily, and blocks until the condition is signaled.

Each time the thread awakens (Line \ref{line:bounded:notfull}),

it checks whether the  
`fPermits` field is positive (that is, the queue is not empty)  
and if not, goes back to sleep.

## Enq Method Part Deux

```
public void enq(T x) {  
    ...  
    if (mustWakeDequuers)  
        deqLock.lock();  
    try {  
        notEmptyCondition.signalAll();  
    } finally {  
        deqLock.unlock();  
    }  
}
```

**Lock and unlock  
deq lock**

Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

79

The `deq()` method is symmetric.

## Enq Method Part Deux

Signal dequeuers that queue no longer empty

```
    deqLock.lock();
    try {
        notEmptyCondition.signalAll();
    } finally {
        deqLock.unlock();
    }
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

80

## The Enq() & Deq() Methods

- Share no locks
  - That's good
- But do share an atomic counter
  - Accessed on every method call
  - That's not so good
- Can we alleviate this bottleneck?

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

81

The deq() method is symmetric.

The key insight is that the enqueuer only decrements the counter, and really cares only whether or not the counter value is zero. Symmetrically, the dequeuer cares only whether or not the counter value is the queue capacity.

NB: the deq() method does not explicitly check the permits field, but relies on testing the sentinel Node's next field. Same thing.

## Split the Counter

- The `enq()` method
  - Decrement only
  - Cares only if value is `zero`
- The `deq()` method
  - Increment only
  - Cares only if value is `capacity`

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

82

The key insight is that the enqueuer only decrements the counter, and really cares only whether or not the counter value is zero. Symmetrically, the dequeuer cares only whether or not the counter value is the queue capacity.

NB: the `deq()` method does not explicitly check the permits field, but relies on testing the sentinel Node's next field. Same thing.

## Split Counter

- Enqueuer decrements `enqSidePermits`
- Dequeuer increments `deqSidePermits`
- When enqueuer runs out
  - Locks `deqLock`
  - Transfers permits
- Intermittent synchronization
  - Not with each method call
  - Need both locks! (careful ...)

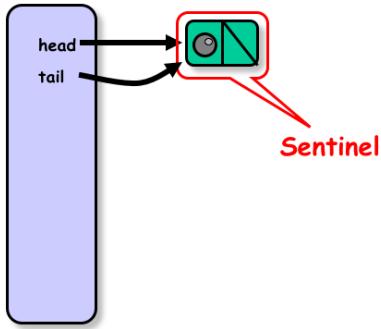
Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

83

Let's summarize. We split the permits field into two parts. The enqueuer decrements the `enqSidePermits` field, and the dequeuer increments the `deqSidePermits` field. When the enqueuer discovers that its field is zero, then it tries to transfer the dequeuer's permits to itself. Why do this? It replaces synchronization with each method call with sporadic, intermittent synchronization. As a practical matter, an enqueuer that runs out of permits needs to acquire the `dequeue` lock (why?) (ANSWER: to prevent the dequeuer from incrementing the value at the same time we are trying to copy it).

WARNING: Any time you try to acquire a lock while holding another, your ears should prick up (alt: your spider-sense should tingle) because you are just asking for a deadlock. So far, no danger, because all of the methods seen so far never hold more than one lock at a time.

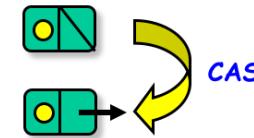
## A Lock-Free Queue



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

84

## Compare and Set



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

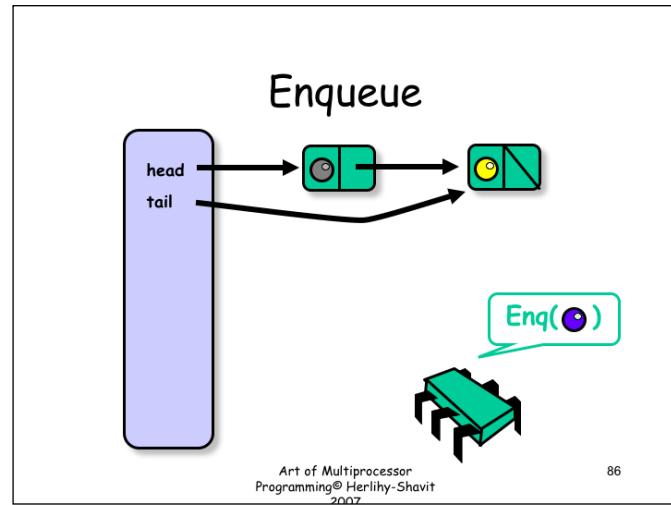
85

Now we consider an alternative approach that does not use locks. The Lock-free queue is due to Maged Michael & Michael Scott

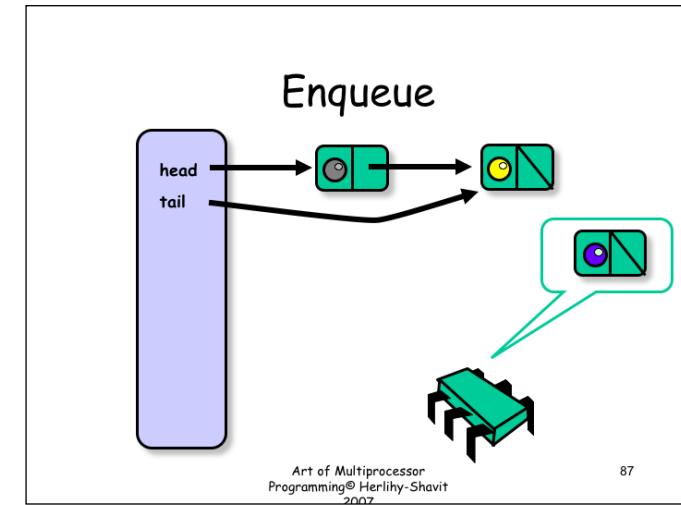
HUMOR: this construction is due to Maged Michael and Michael Scott. Everyone cites it as Michael Scott, which means that many people (unjustly) think that Michael Scott invented it. The moral is, make sure that your last name is not the same as your advisor's first name.

Here too we use a list-based structure. Since the queue is unbounded, arrays would be awkward. As before, We start out with head and tail fields that point to the first and last entries in the list. The first Node in the list is a sentinel Node whose value field is meaningless. The sentinel acts as a placeholder.

One difference from the lock-based queue is that we use CAS to manipulate the next fields of entries. Before, we did all our synchronization on designed fields in the queue itself, but here we must synchronize on Node fields as well.

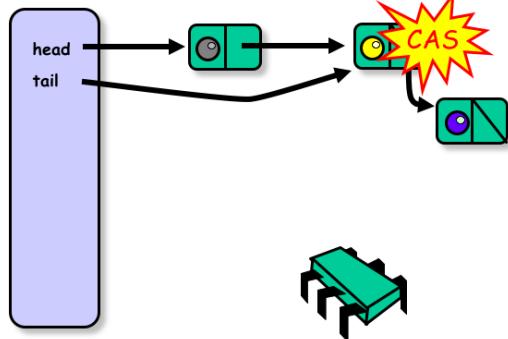


As a first step, the enqueuer calls CAS to sent the tail Node's next field to the new Node.



As a first step, the enqueuer calls CAS to sent the tail Node's next field to the new Node.

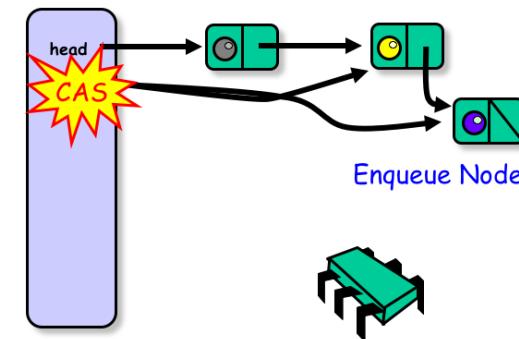
## Logical Enqueue



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

88

## Physical Enqueue



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

89

As a first step, the enqueuer calls CAS to sent the tail Node's next field to the new Node.

Once the tail Node's next field has been redirected, use CAS to redirect the queue's tail field to the new Node.

## Enqueue

- These two steps are not atomic
- The tail field refers to either
  - Actual last Node (good)
  - Penultimate Node (not so good)
- Be prepared!

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

90

[pe·nul·ti·mate - adjective -- next to the last: "the penultimate scene of the play"]

As you may have guessed, the situation is more complicated than it appears. The enq() method is supposed to be atomic, so how can we get away with implementing it in two non-atomic steps? We can count on the following invariant: the tail field is a reference to (1) actual last Node, or (2) the Node just before the actual last Node.

## Enqueue

- What do you do if you find
  - A trailing tail?
- Stop and help fix it
  - If tail node has non-null next field
  - CAS the queue's tail field to tail.next
- As in the universal construction

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

91

Any method call must be prepared to encounter a tail field that is trailing behind the actual tail. This situation is easy to detect. How? (Answer: tail.next not null). How to deal with it? Fix it! Call compareAndSet to set the queue's tail field to point to the actual last Node in the queue!

## When CASs Fail

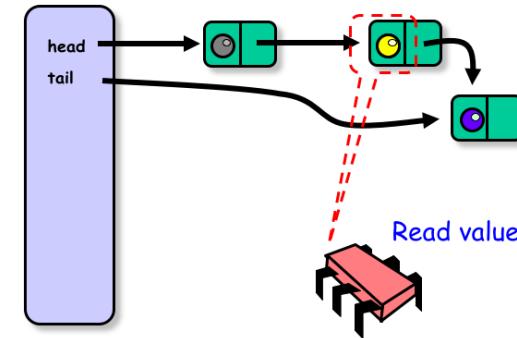
- During logical enqueue
  - Abandon hope, restart
  - Still lock-free (why?)
- During physical enqueue
  - Ignore it (why?)

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

92

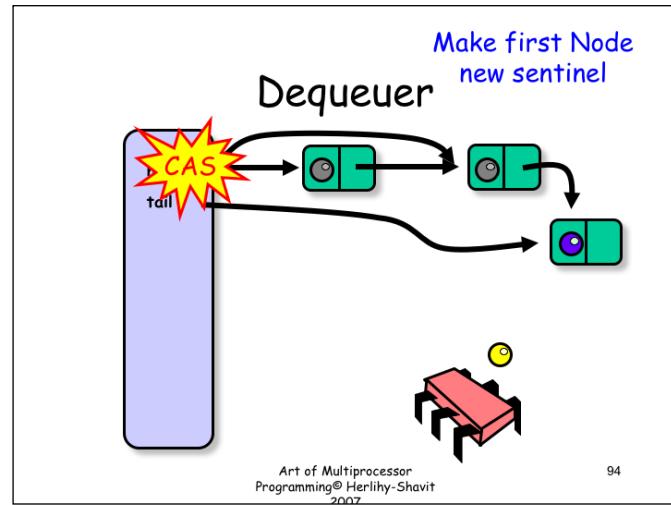
What do we do if the CAS calls fail? It matters a lot in which CAS call this happened. In Step One, a failed CAS implies a synchronization conflict with another enq() call. Here, the most sensible approach is to abandon the current effort (panic!) and start over. In Step Two, we can ignore a failed CAS, simply because the failure means that some other thread has executed that step on our behalf.

## Dequeueer



93

Next, the thread reads the first Node's value, and stores it in a local variable. The slide suggests that the value is nulled out, but in fact there is no need to do so.



The thread then makes the first Node the new sentinel, and discards the old sentinel. This is the clever trick I mentioned earlier that ensures that we do not need to lock the Node itself. We do not physically remove the Node, we just demote it to sentinel.

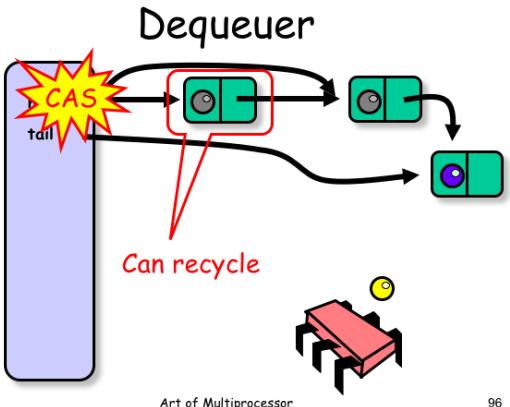
## Memory Reuse?

- What do we do with nodes after we dequeue them?
- Java: let garbage collector deal?
- Suppose there is no GC, or we prefer not to use it?

95

Art of Multiprocessor Programming® Herlihy-Shavit 2007

Let's take a brief detour. What do we do with nodes after we remove them from the queue? This is not an issue with the lock-based implementation, but it requires some thought here. In Java, we can just let the garbage collector recycle unused objects. But what if we are operating in an environment where there is no garbage collector, or where we think we can recycle memory more efficiently?



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

96

Let us look at the dequeue method from a different perspective. When we promote the prior first Node to sentinel, what do we do with the old sentinel? Seems like a good candidate for recycling.

## Simple Solution

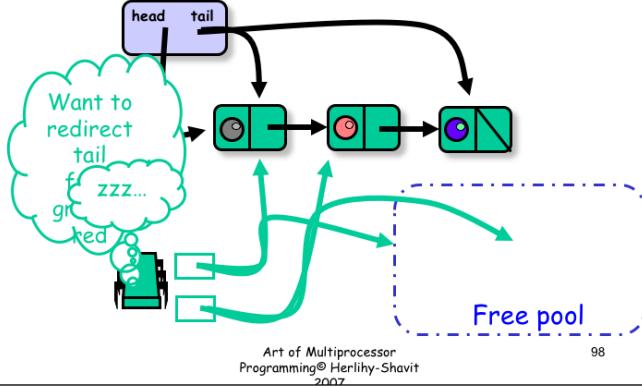
- Each thread has a free list of unused queue nodes
- Allocate node: pop from list
- Free node: push onto list
- Deal with underflow somehow ...

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

97

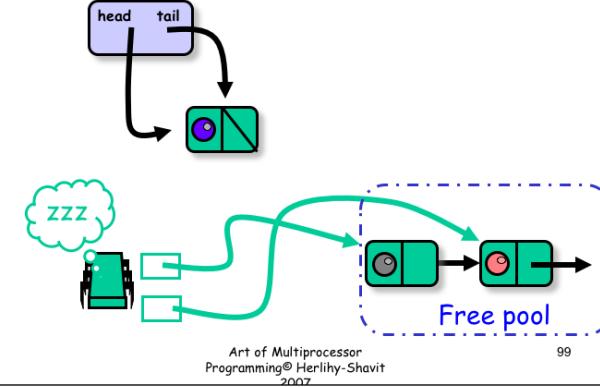
The most reasonable solution is to have each thread manage its own pool of unused nodes. When a thread needs a new Node, it pops one from the list. No need for synchronization. When a thread frees A Node, it pushes the newly-freed Node onto the list. What do we do when a thread runs out of nodes? Perhaps we could just malloc() more, or perhaps we can devise a shared pool. Never mind.

## Why Recycling is Hard



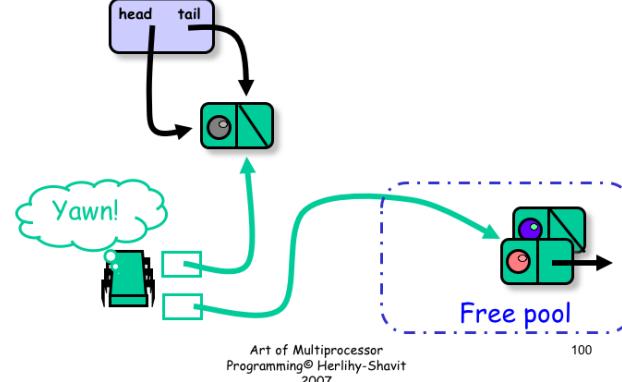
Now the green thread goes to sleep, and other threads dequeue the red object and the green object, sending the prior sentinel and prior red Node to the respective free pools of the dequeuing threads.

## Both Nodes Reclaimed



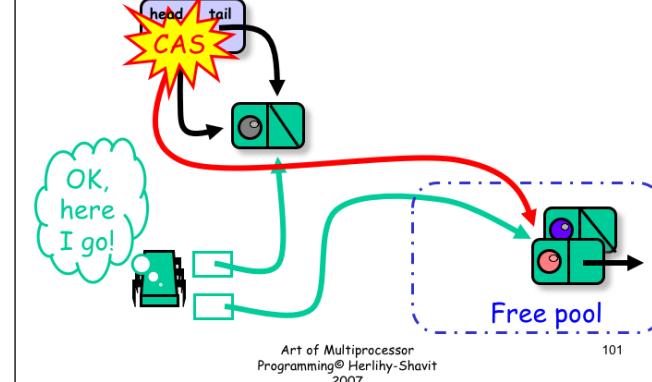
Despite what you might think, we are perfectly safe, because any thread that tries to CAS the head field must fail, because that field has changed. Now assume that the enough `deq()` and `enq()` calls occur such that the original sentinel Node is recycled, and again becomes a sentinel Node for an empty queue.

## One Node Recycled

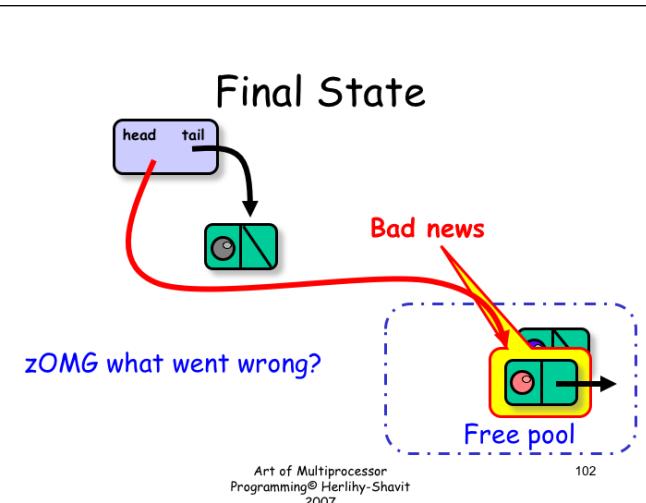


Now the green thread wakes up. It applies *CAS* to the queue's head field ..

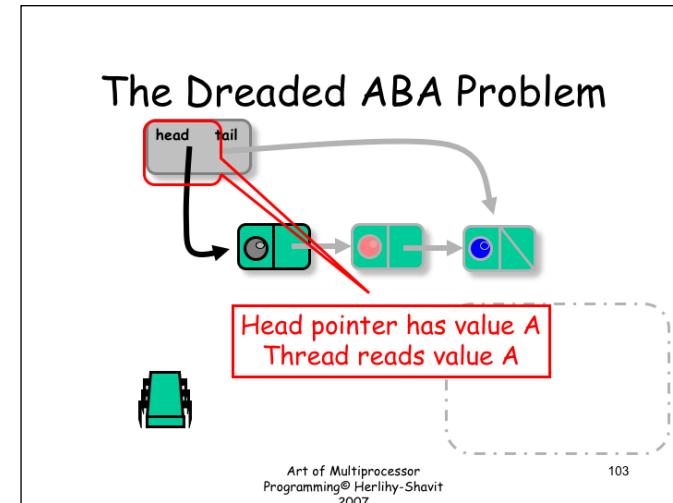
## Why Recycling is Hard



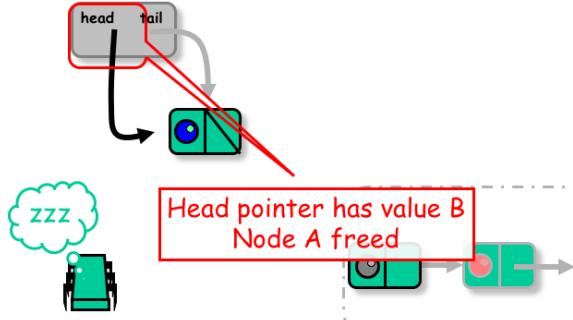
Surprise! It works! The problem is that the bit-wise value of the head field is the same as before, even though its meaning has changed. This is a problem with the way the *CAS* operation is defined, nothing more.



In the end, the tail pointer points to a sentinel node, while the head points to node in some thread's free list. What went wrong?



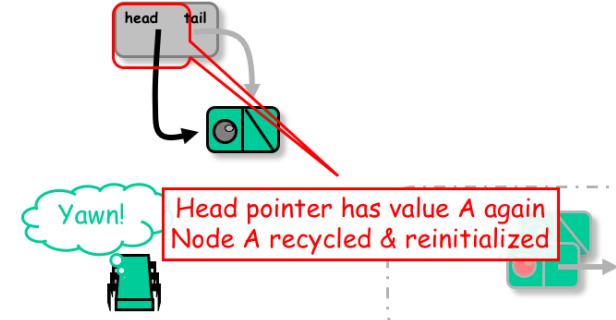
## Dreaded ABA continued



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

104

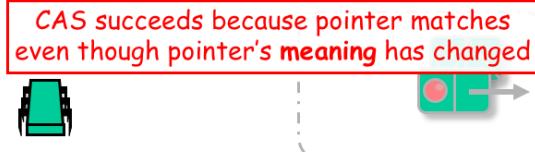
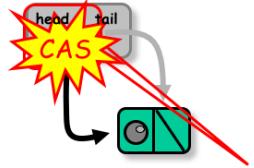
## Dreaded ABA continued



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

105

## Dreaded ABA continued



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

106

## The Dreaded ABA Problem

- Is a result of CAS() semantics
  - I blame Sun, Intel, AMD, ...
- Not with Load-Locked/Store-Conditional
  - Good for IBM?

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

107

## Dreaded ABA - A Solution

- Tag each pointer with a counter
- Unique over lifetime of node
- Pointer size vs word size issues
- Overflow?
  - Don't worry be happy?
  - Bounded tags?
- AtomicStampedReference class

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

108

## Atomic Stamped Reference

- AtomicStampedReference class
  - Java.util.concurrent.atomic package

Can get reference and stamp atomically



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

109

The actual implementation in Java is by having an indirection through a special sentinel node. If the special node exists then the mark is set, and if the node does not, then the bit is false. The sentinel points to the desired address so we pay a level of indirection. In C, using alignment on 32 bit architectures or on 64 bit architectures, we can "steal" a bit from the pointer.

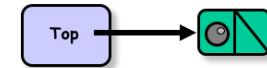
## Concurrent Stack

- Methods
  - `push(x)`
  - `pop()`
- Last-in, First-out (LIFO) order
- Lock-Free!

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

110

## Empty Stack



Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

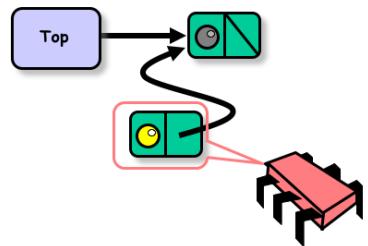
111

To `\mPush{}` an item of type `\cT{}` into the stack  
(As in our `\cLockFreeQueue{}` implementation, we do not allow  
insertion of a `\Jnull{}` value), a  
thread creates a new node (Line~`\ref{line:elimination:new}`),  
and then calls the `\mTryPush{}` method that directs its next pointer to  
the node  
pointed-to by the `\fTop{}`, and then attempts to swing the `\fTop{}`  
to point to the new node using a `\mCompareAndSet{}` method call. If  
`\mTryPush{}` succeeds,  
`\mPush{}` returns, and if not, the `\mTryPush{}` attempt is repeated.

110

111

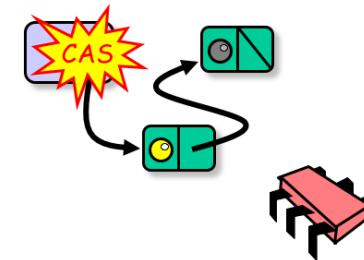
## Push



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

112

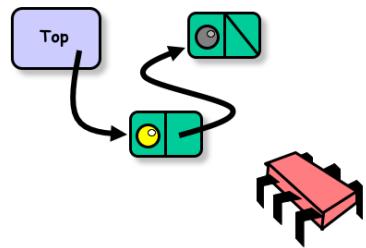
## Push



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

113

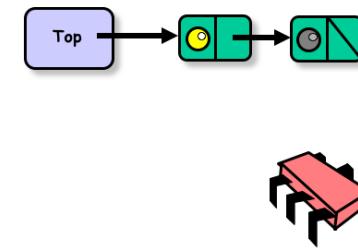
## Push



Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

114

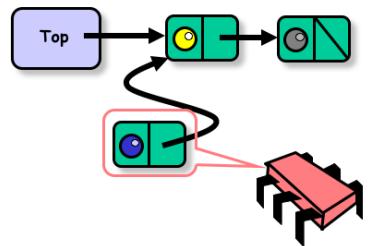
## Push



Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

115

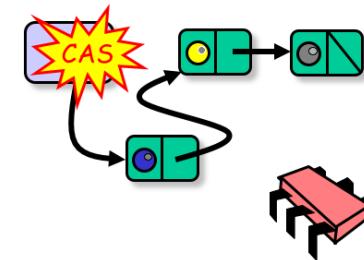
## Push



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

116

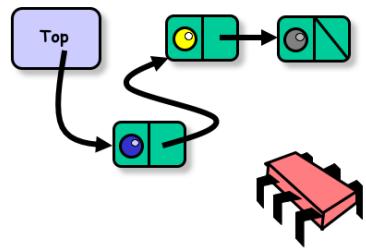
## Push



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

117

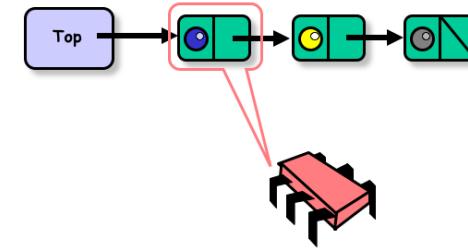
## Push



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

118

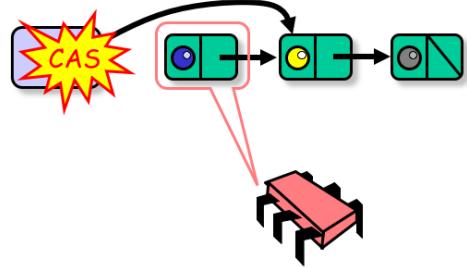
## Pop



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

119

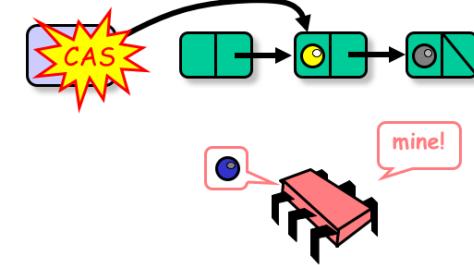
Pop



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

120

Pop



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

121

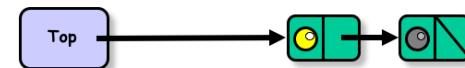
Pop



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

122

Pop



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

123

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top =  
        new AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node)) {  
                return;  
            } else backoff.backoff();  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

124

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
        AtomicReference(null);  
    public Boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node))  
                return;  
        }  
    }  
}
```

tryPush attempts to push a node

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

125

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
    AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node)) {  
                return; Read top value  
            } else backoff.backoff()  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

126

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
    AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node)) {  
                return;  
            } else backoff.backoff()  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

127

**current top will be new node's successor**

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
    AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node)) {  
                return:  
            }  
        }  
    }  
}
```

Try to swing top, return success or failure

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

128

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
    AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node)) {  
                return;  
            } else backoff  
        }  
    }  
}
```

Push calls tryPush

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

129

## Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
    AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldTop = top.get();  
        node.next = oldTop;  
        return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node))  
                return;  
            } else backoff.backoff()  
    }}  
  
Create new node
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

130

## Lock-free Stack

Makes scheduling benevolent so  
Method is effectively wait-  
free

```
public class LockFreeStack {  
    private AtomicReference top = new  
    AtomicReference(null);  
    public boolean tryPush(Node node){  
        Node oldT  
        node.next  
        return(to  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node)) {  
                return;  
            } else backoff.backoff()  
    }}  
  
If tryPush() fails,  
back off before retrying
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

131

## Lock-free Stack

- Good
  - No locking
- Bad
  - Without GC, fear ABA
  - Without backoff, huge contention at top
  - In any case, no parallelism

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

132

## Big Question

- Are stacks *inherently sequential*?
- Reasons why
  - Every `pop()` call fights for top item
- Reasons why not
  - Stay tuned ...

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

133

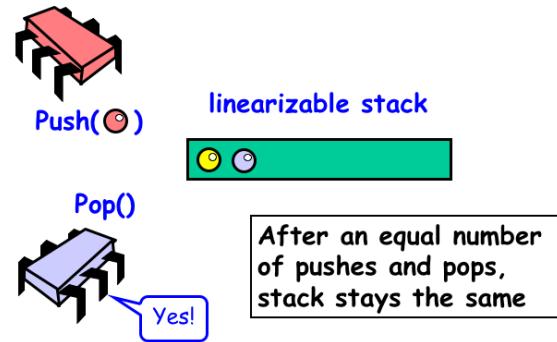
## Elimination-Backoff Stack

- How to
  - "turn contention into parallelism"
- Replace familiar
  - exponential backoff
- With alternative
  - elimination-backoff

Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

134

## Observation

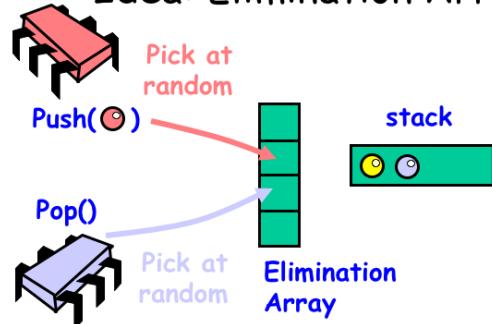


Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

135

Take any linearizable stack, for example, the lock-free queue.

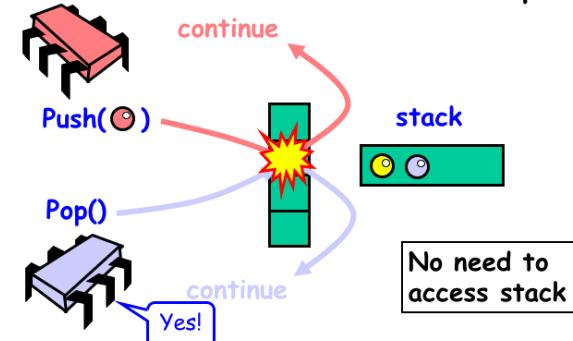
### Idea: Elimination Array



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

136

### Push Collides With Pop



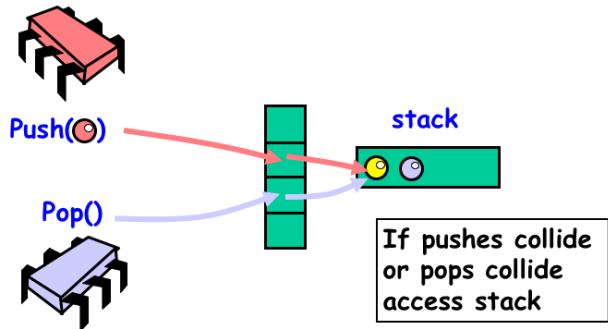
Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

137

Pick random locations in the array.

No need to access stack since it stays the same anyhow. Just exchange the values.

## No Collision



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

138

## Elimination-Backoff Stack

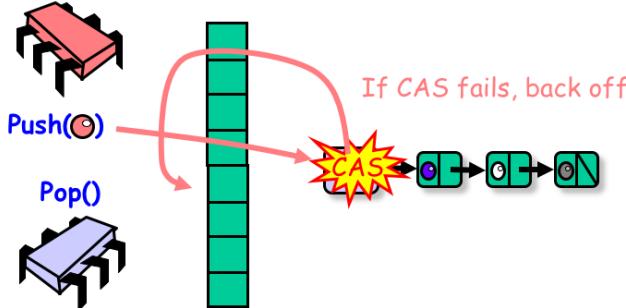
- Lock-free stack + elimination array
- Access Lock-free stack,
  - If uncontended, apply operation
  - if contended, back off to elimination array and attempt elimination

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

139

No need to access stack since it stays the same anyhow. Just exchange the values.

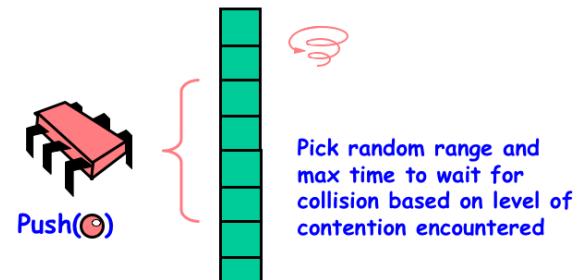
## Elimination-Backoff Stack



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

140

## Dynamic Range and Delay



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

141

No need to access stack since it stays the same anyhow. Just exchange the values.

As the contention grows,

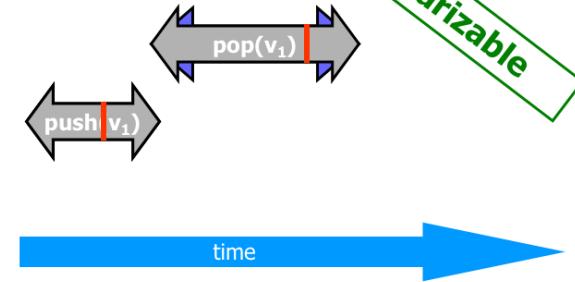
## Linearizability

- Un-eliminated calls
  - linearized as before
- Eliminated calls:
  - linearize pop() immediately after matching push()
- Combination is a linearizable stack

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

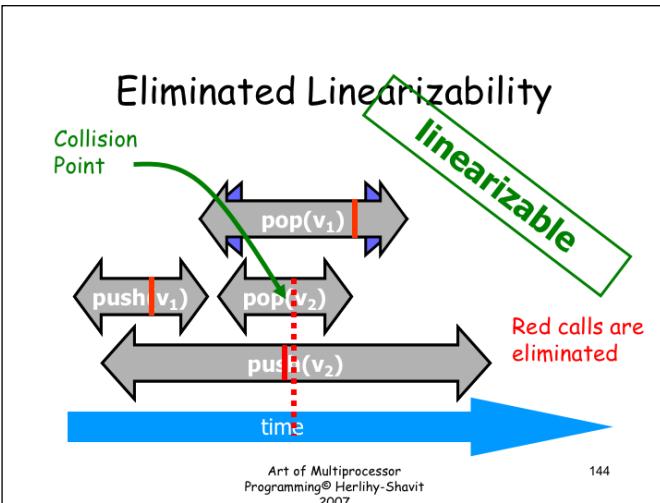
142

## Un-Eliminated Linearizability



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

143



## Backoff Has Dual Effect

- Elimination introduces parallelism
- Backoff onto array cuts contention on lock-free stack
- Elimination in array cuts down total number of threads ever accessing lock-free stack

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

145

## Elimination Array

```
public class EliminationArray {  
    private static final int duration = ...;  
    private static final int timeUnit = ...;  
    Exchanger<T>[] exchanger;  
    public EliminationArray(int capacity) {  
        exchanger = new Exchanger[capacity];  
        for (int i = 0; i < capacity; i++)  
            exchanger[i] = new Exchanger<T>();  
        ...  
    }  
    ...  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

146

## Elimination Array

```
public class EliminationArray {  
    private static final int duration = ...;  
    private static final int timeUnit = ...;  
    Exchanger<T>[] exchanger;  
    public EliminationArray(int capacity) {  
        exchanger = new Exchanger[capacity];  
        for (int i = 0; i < capacity; i++)  
            exchanger[i] = new Exchanger<T>();  
        ...  
    }  
    ...  
}
```

An array of exchangers

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

147

## A Lock-Free Exchanger

```
public class Exchanger<T> {  
    AtomicStampedReference<T> slot  
        = new AtomicStampedReference<T>(null, 0);
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

148

## A Lock-Free Exchanger

```
public class Exchanger<T> {  
    AtomicStampedReference<T> slot  
        = new AtomicStampedReference<T>(null, 0);
```

Atomically modifiable  
reference + time stamp

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

149

## Atomic Stamped Reference

- **AtomicStampedReference class**
  - Java.util.concurrent.atomic package
- In C or C++:



Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

150

## Extracting Reference & Stamp

```
Public T get(int[] stampHolder);
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

151

The actual implementation in Java is by having an indirection through a special sentinel node. If the special node exists then the mark is set, and if the node does not, then the bit is false. The sentinel points to the desired address so we pay a level of indirection. In C, using alignment on 32 bit architectures or on 64 bit architectures, we can "steal" a bit from the pointer.

## Extracting Reference & Stamp

```
Public T get(int[] stampHolder);
```

Returns reference to object of type T

Returns stamp at array index 0!

Art of Multiprocessor Programming © Herlihy-Shavit 2007

152

## Exchanger Status

```
enum Status {EMPTY, WAITING, BUSY};
```

Art of Multiprocessor Programming © Herlihy-Shavit 2007

153

## Exchanger Status

```
enum Status {EMPTY, WAITING, BUSY};
```

Nothing yet

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

154

## Exchange Status

```
enum Status {EMPTY, WAITING, BUSY};
```

Nothing yet

One thread is waiting  
for rendez-vous

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

155

## Exchange Status

```
enum Status {EMPTY, WAITING, BUSY};
```

Nothing yet

One thread is waiting  
for rendez-vous

Other threads busy  
with rendez-vous

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

156

## The Exchange

```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: ... // slot is free
            case WAITING: ... // someone waiting for me
            case BUSY: ... // others exchanging
        }
    }
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

157

## The Exchange

```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: ... // slot is free
            case WAITING: ... // someone waiting for me
            case BUSY: ... // others exchanging
        }
    }
}
```

Item & timeout

158

## The Exchange

```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true)
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: ... // slot is free
            case WAITING: ... // someone waiting for me
            case BUSY: ... // others exchanging
        }
}
```

Array to hold timestamp

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

159

## The Exchange

```
public T Exchange(T myItem, long nanos) throws  
TimeoutException {  
    long timeBound = System.nanoTime() + nanos;  
    int[] stampHolder = {0};  
    while (true) {  
        if (System.nanoTime() > timeBound)  
            throw new TimeoutException();  
        T herItem = slot.get(stampHolder);  
        int stamp = stampHolder[0];  
        switch(stamp)  
            case EMPTY: // slot is free  
            case WAITING: // someone waiting for me  
            case BUSY: // others exchanging  
        }  
    } } Loop until timeout
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

160

## The Exchange

```
public T Exchange(T myItem, long nanos) throws  
TimeoutException {  
    long timeBound = System.nanoTime() + nanos;  
    int[] stampHolder = {0};  
    while (true) {  
        if (System.nanoTime() > timeBound)  
            throw new TimeoutException();  
        T herItem = slot.get(stampHolder);  
        int stamp = stampHolder[0];  
        switch(stamp)  
            case EMPTY: // slot is free  
            case WAITING: // someone waiting for me  
            case BUSY: // others exchanging  
        } } Get other's item and timestamp
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

161

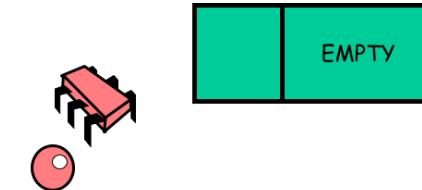
## The Exchange

```
public T Exchange(T myItem, long nanos) throws  
TimeoutException {  
    long timeBound = System.nanoTime() + nanos;  
Exchanger slot has three states  
    while (true) {  
        if (System.nanoTime() > timeBound)  
            throw new TimeoutException();  
        T herItem = slot.get(stampHolder);  
        int stamp = stampHolder[0];  
        switch(stamp) {  
            case EMPTY: ... // slot is free  
            case WAITING: ... // someone waiting for me  
            case BUSY: ... // others exchanging  
        }  
    }}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

162

## Lock-free Exchanger



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

163

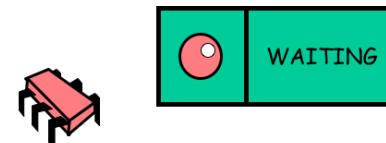
## Lock-free Exchanger



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

164

## Lock-free Exchanger



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

165

## Lock-free Exchanger

In search of partner ...



Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

166

## Lock-free Exchanger

Still waiting ...



Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

167

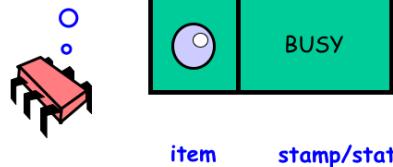
166

167

Obviously some other thread can try and set the state to 2 also and only the winner does the exchange.

## Lock-free Exchanger

Partner showed  
up, take item and  
reset to EMPTY

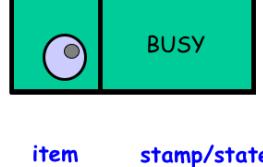
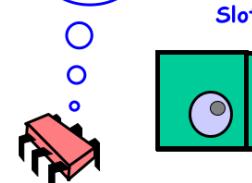


Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

168

## Lock-free Exchanger

Partner showed  
up, take item and  
reset to EMPTY



Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

169

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, EMPTY,
    WAITING)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
            if (slot.compareAndSet(myItem, null, WAITING,
            EMPTY)){throw new TimeoutException();}
            else {
                herItem = slot.get(stampHolder);
                slot.set(null, EMPTY);
                return herItem;
            }
        } break;
    }
```

Programming - Remy Saviet  
2007

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, EMPTY,
    WAITING)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
            if (slot.compareAndSet(myItem, null, WAITING,
            EMPTY)){throw new TimeoutException();}
            else {
                herItem = slot.get(stampHolder);
                slot.set(null, EMPTY);
                return herItem;
            }
        } break;
    }
```

Slot is free, try to insert  
myItem and change state  
to WAITING

Programming - Remy Saviet  
2007

We present algorithm with timestamps so that you will not worry about possible ABA but actually can implement this scheme with counter that has just 3 values 0,1, and 2.

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, EMPTY,
    WAITING)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
            if (slot.compareAndSet(myItem, null, WAITING,
            EMPTY)){throw new TimeoutException();}
        } else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    } break;
```

Programming - Jeremy Shaw  
2007

Loop while still time left to  
attempt exchange

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, WAITING,
    BUSY)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
            if (slot.compareAndSet(myItem, null, WAITING,
            EMPTY)){throw new TimeoutException();}
        } else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    } break;
```

Programming - Jeremy Shaw  
2007

Get item and stamp in slot  
and check if state changed  
to BUSY

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, EMPTY,
    WAITING)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
            if (slot.compareAndSet(myItem, null, WAITING,
            EMPTY)){throw new TimeoutException();
            } else {
                herItem = slot.get(stampHolder);
                slot.set(null, EMPTY);
                return herItem;
            }
        } break;
    }
```

Programming - Jeremy Shaw  
2007

If successful reset slot  
state to EMPTY

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, WAITING,
    BUSY)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
            if (slot.compareAndSet(myItem, null, WAITING,
            EMPTY)){throw new TimeoutException();
            } else {
                herItem = slot.get(stampHolder);
                slot.set(null, EMPTY);
                return herItem;
            }
        } break;
    }
```

Programming - Jeremy Shaw  
2007

and return item found in slot

Notice that we do not need a CAS here, only a set which is a write, because only the thread that changed the state from 0 to 1 can reset it from 2 to 0.

## Exchanger State EMPTY

Otherwise we ran out of time, try to reset state to EMPTY, if successful time out

```
if (stampHolder[0] == BUSY) {
    slot.set(null, EMPTY);
    return herItem;
}
if (slot.compareAndSet(myItem, null, WAITING,
EMPTY)){throw new TimeoutException();
} else {
    herItem = slot.get(stampHolder);
    slot.set(null, EMPTY);
    return herItem;
}
} break;
```

## Exchanger State EMPTY

If reset failed, someone showed up after all, take that item

```
while (System.nanoTime() < timeBound){
    herItem = slot.get(stampHolder);
    if (stampHolder[0] == BUSY) {
        slot.set(null, EMPTY);
        return herItem;
    }
    if (slot.compareAndSet(myItem, null, WAITING,
EMPTY)){throw new TimeoutException();
} else {
    herItem = slot.get(stampHolder);
    slot.set(null, EMPTY);
    return herItem;
}
} break;
```

## Exchanger State EMPTY

```
Set slot to EMPTY with new
timestamp and return item found
    while (System.nanoTime() < timeBound){
        herItem = slot.get(stampHolder);
        if (stampHolder[0] == BUSY) {
            slot.set(null, EMPTY);
            return herItem;
        }
        if (slot.compareAndSet(myItem, null, WAITING,
EMPTY)){throw new TimeOutException();}
        else{
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    } break;
```

Again notice that we do not need a CAS here, only a set which is a write, because only the thread that changed the state from 0 to 1 can reset it from 2 to 0...

## Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, EMPTY,
WAITING)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY)
                slot.set(null, EMPTY);
            return herItem;
        }
        if (slot.compareAndSet(myItem, null, WAITING,
EMPTY)){throw new TimeOutException();}
        else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    } break;
```

Again notice that we do not need a CAS here, only a set which is a write, because only the thread that changed the state from 0 to 1 can reset it from 2 to 0...

## States WAITING and BUSY

```
case WAITING: // someone waiting for me
    if (slot.compareAndSet(herItem, myItem,
        WAITING, BUSY))
        return herItem;
    break;
case BUSY: // others in middle of exchanging
    break;
default: // impossible
    break;
}
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

180

## States WAITING and BUSY

```
case WAITING: // someone waiting for me
    if (slot.compareAndSet(herItem, myItem,
        WAITING, BUSY))
        return herItem;
    break;
case BUSY: // others in middle of exchanging
    break;
default: state WAITING means
    someone is waiting for an
    exchange, so attempt to
    CAS my item in and change
    state to BUSY
}
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

181

### States WAITING and BUSY

```
case WAITING: // someone waiting for me
    if (slot.compareAndSet(herItem, myItem,
                           WAITING, BUSY))
        return herItem;
    break;

case BUSY: // others in middle of exchanging
    break;
default: // impossible
    break;
}
```

If successful return her item, state is now BUSY, otherwise someone else took her item so try again from start

Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

182

## States WAITING and BUSY

```
case WAITING: // someone waiting for me
    if (slot.compareAndSet(herItem, myItem,
WAITING, BUSY))
        return herItem;
    break;

case BUSY: // others in middle of exchanging
break;

default: // impossible
    break;
}
}
}
```

If state is BUSY then  
some other threads are  
using slot to exchange so  
start again

Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

18

## The Exchanger Slot

- Exchanger is lock-free
- Because the only way an exchange can fail is if others repeatedly succeeded or no-one showed up
- The slot we need does not require symmetric exchange

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

184

## Elimination Array

```
public class EliminationArray {  
    ...  
    public T visit(T value, int Range) throws  
        TimeoutException {  
        int slot = random.nextInt(Range);  
        int nanodur = convertToNanos(duration, timeUnit));  
        return (exchanger[slot].exchange(value, nanodur)  
    }  
}
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

185

Notice that The slot we need does not require symmetric exchange, only  
push needs to pass item to pop.

## Elimination Array

```
public class EliminationArray {  
    public T visit(T value, int Range) throws  
        TimeoutException {  
        int slot = random.nextInt(Range);  
        int nanodur = convertToNanos(duration, timeUnit);  
        return (exchanger[slot].exchange(value, nanodur)  
    }}
```

visit the elimination array with a value  
and a range (duration to wait is not  
dynamic)

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

186

## Elimination Array

```
public class EliminationArray {  
    ...  
    public T visit(T value, int Range) throws  
        TimeoutException {  
        Pick a random array entry  
        int slot = random.nextInt(Range);  
        int nanodur = convertToNanos(duration, timeUnit));  
        return (exchanger[slot].exchange(value, nanodur)  
    }}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

187

## Elimination Array

```
public class EliminationArray {  
    ...  
    public T Exchange value or time out  
    TimeoutException t  
    int slot = random.nextIntRange();  
    int nanodur = convertToNanos(duration, timeUnit);  
    return (exchanger[slot].exchange(value, nanodur))  
}
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

188

## Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
                eliminationArray.visit(value,policy.Range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Art of Multiprocessor  
Programming© Herlihy-Shavit  
2007

189

## Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
                eliminationArray.visit(value,policy.Range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

First try to push

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

190

## Elimination Stack Push

```
public void push(T value) {  
    ...  
    whi If failed back-off to try to eliminate  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
                eliminationArray.visit(value,policy.Range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

191

## Elimination Stack Push

```
public void push(T value) {  
    ...  
    while Value being pushed and range to try {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
                eliminationArray.visit(value,policy.Range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

192

## Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (Only a pop has null value  
so elimination was successful) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
                eliminationArray.visit(value,policy.Range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

193

## Elimination Stack Push

```
public void push(T value) {  
    ...  
    Else retry push on lock-free stack  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
                eliminationArray.visit(value,policy.Range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

194

## Elimination Stack Pop

```
public T pop() {  
    ...  
    while (true) {  
        if (tryPop()) {  
            return returnNode.value;  
        } else  
            try {  
                T otherValue =  
                    eliminationArray.visit(null,policy.Range);  
                if (otherValue != null) {  
                    return otherValue;  
                }  
            }  
    }  
}
```

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

195

## Elimination Stack Pop

```
public T popO {  
    ...  
    while (...) {  
        if (...) {  
            re...  
        } else {  
            try {  
                T otherValue =  
                    eliminationArray.visit(null, policy.Range);  
                if (otherValue != null) {  
                    return otherValue;  
                }  
            }  
        }  
    }}}
```

Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

196

Same as push, if non-null other  
thread must have pushed  
so elimination succeeds

## Summary

- We saw both lock-based and lock-free implementations of
- queues and stacks
- Don't be quick to declare a data structure inherently sequential
  - Linearizable stack is not inherently sequential
- ABA is a real problem, pay attention

Art of Multiprocessor  
Programming © Herlihy-Shavit  
2007

197



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](#).

- You are free:
  - to Share — to copy, distribute and transmit the work
  - to Remix — to adapt the work
- Under the following conditions:
  - **Attribution.** You must attribute the work to "The Art of Multiprocessor Programming" (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

Art of Multiprocessor  
Programming® Herlihy-Shavit  
2007

198